



---

## *Chap. 13) I/O Systems*

경희대학교 컴퓨터공학과

방재훈

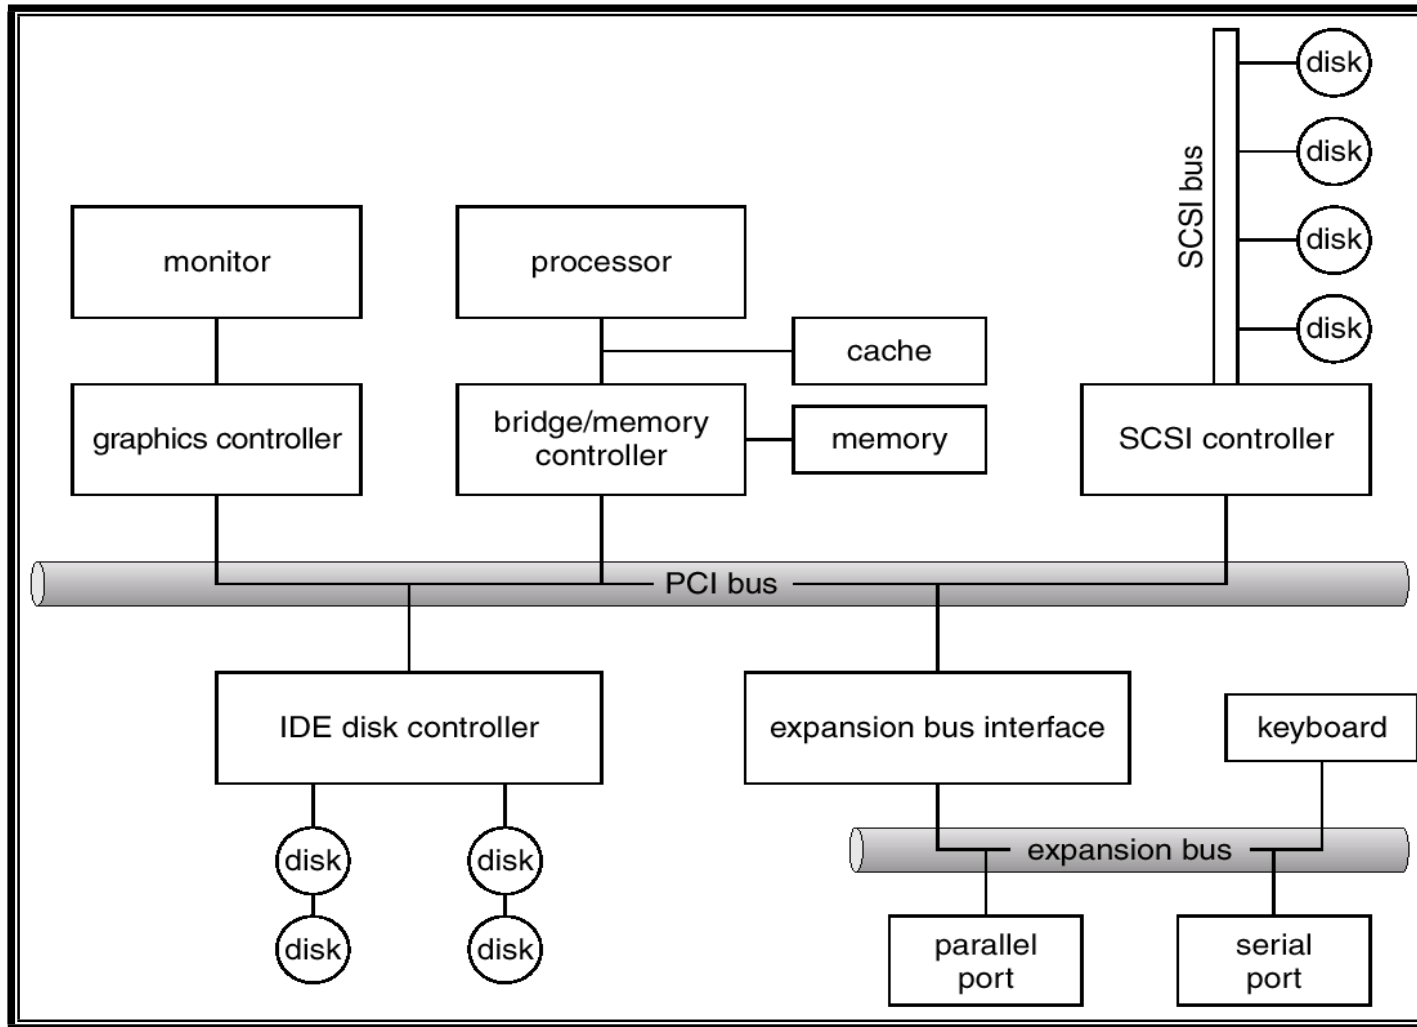
# I/O Hardware

---

- Incredible variety of I/O devices
- Common concepts
  - ✓ Port
  - ✓ Bus (daisy chain or shared direct access)
  - ✓ Controller (host adapter)
- I/O instructions control devices
- Devices have addresses, used by
  - ✓ Direct I/O instructions
  - ✓ Memory-mapped I/O



# A Typical PC Bus Structure



## ■ Device controller (or host adapter)

- ✓ I/O devices have components:
  - Mechanical component
  - Electronic component
- ✓ The electronic component is the device controller
  - May be able to handle multiple devices
- ✓ Controller's tasks
  - Convert serial bit stream to block of bytes
  - Perform error correction as necessary
  - Make available to main memory



- Use special I/O instructions to an I/O port address

I/O address range (hexadecimal)	device
000-00F	DMA controller
020-021	interrupt controller
040-043	timer
200-20F	game controller
2F8-2FF	serial port (secondary)
320-32F	hard-disk controller
378-37F	parallel port
3D0-3DF	graphics controller
3F0-3F7	diskette-drive controller
3F8-3FF	serial port (primary)



- The device control registers are mapped into the address space of the processor
  - ✓ The CPU executes I/O requests using the standard data transfer instructions
- I/O device drivers can be written entirely in C
- No special protection mechanism is needed to keep user processes from performing I/O
  - ✓ Can give a user control over specific devices but not others by simply including the desired pages in its page table
- Reading a device register and testing its value is done with a single instruction
- Memory-mapped regions should be uncacheable
- Memory-mapped device register is vulnerable to accidental modification through the use of incorrect pointers
  - ✓ Protected memory helps to reduce this risk



## ■ Polled I/O

- ✓ CPU asks (“polls”) devices if need attention
  - ready to receive a command
  - command status, etc.
- ✓ Advantages
  - Simple
  - Software is in control
  - Efficient if CPU finds a device to be ready soon
- ✓ Disadvantages
  - Inefficient in non-trivial system (high CPU utilization)
  - Low priority devices may never be serviced



# Polling vs. Interrupts (Cont'd)

---

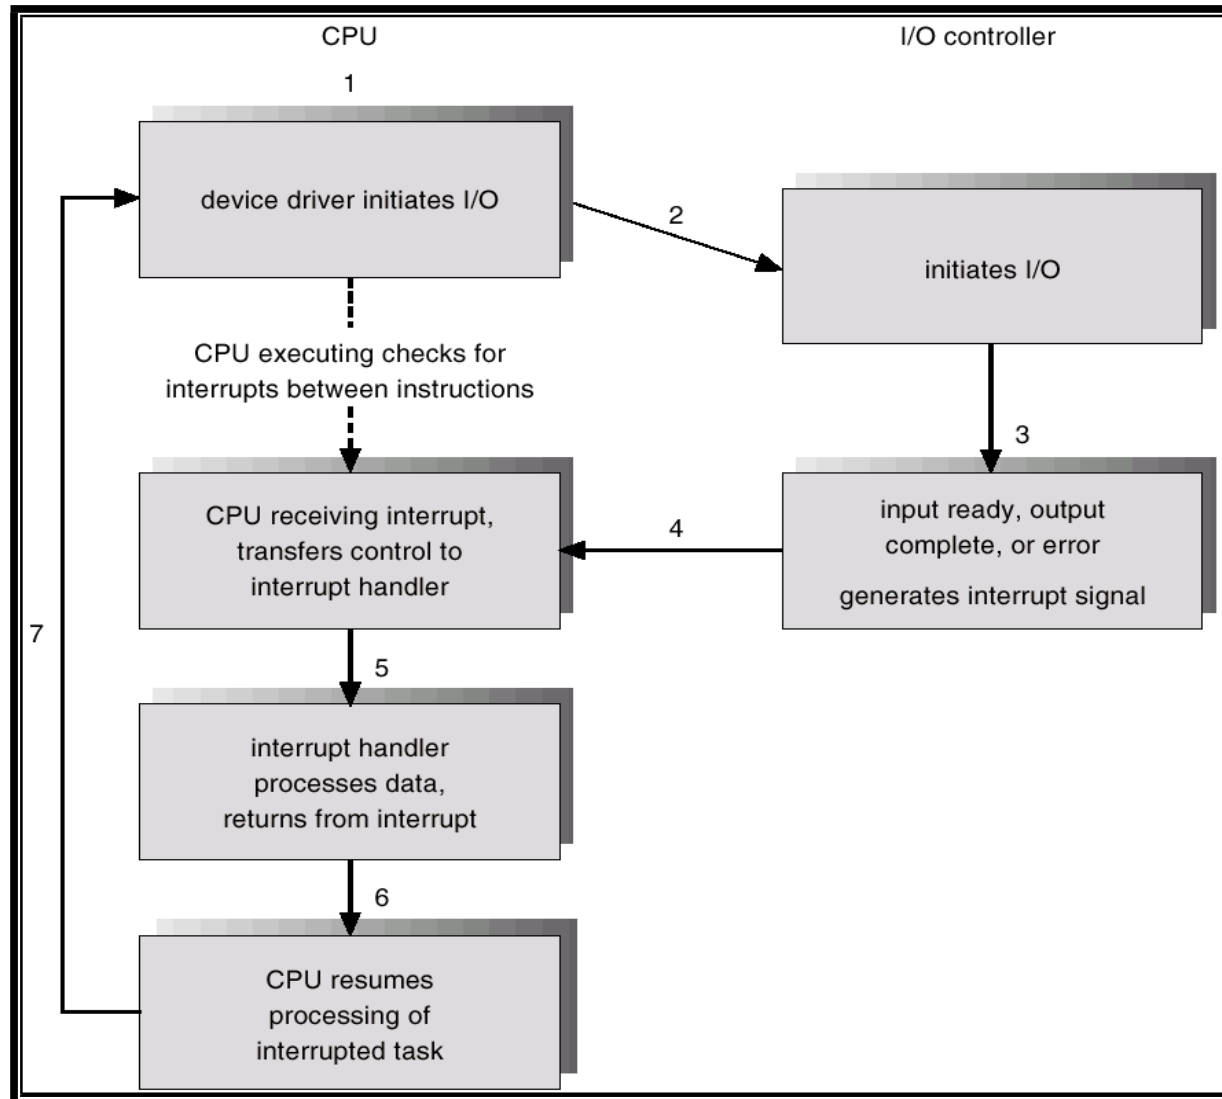
## ■ Interrupt-driven I/O

- ✓ I/O devices request interrupt when need attention
- ✓ Interrupt service routines specific to each device are invoked
- ✓ Interrupts can be shared between multiple devices
- ✓ Advantages
  - CPU only attends to device when necessary
  - More efficient than polling in general
- ✓ Disadvantages
  - Excess interrupts slow (or prevent) program execution
  - Overheads (may need 1 interrupt per byte transferred)





# Interrupt-Driven I/O Cycle



# Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19Ð31	(Intel reserved, do not use)
32Ð255	maskable interrupts



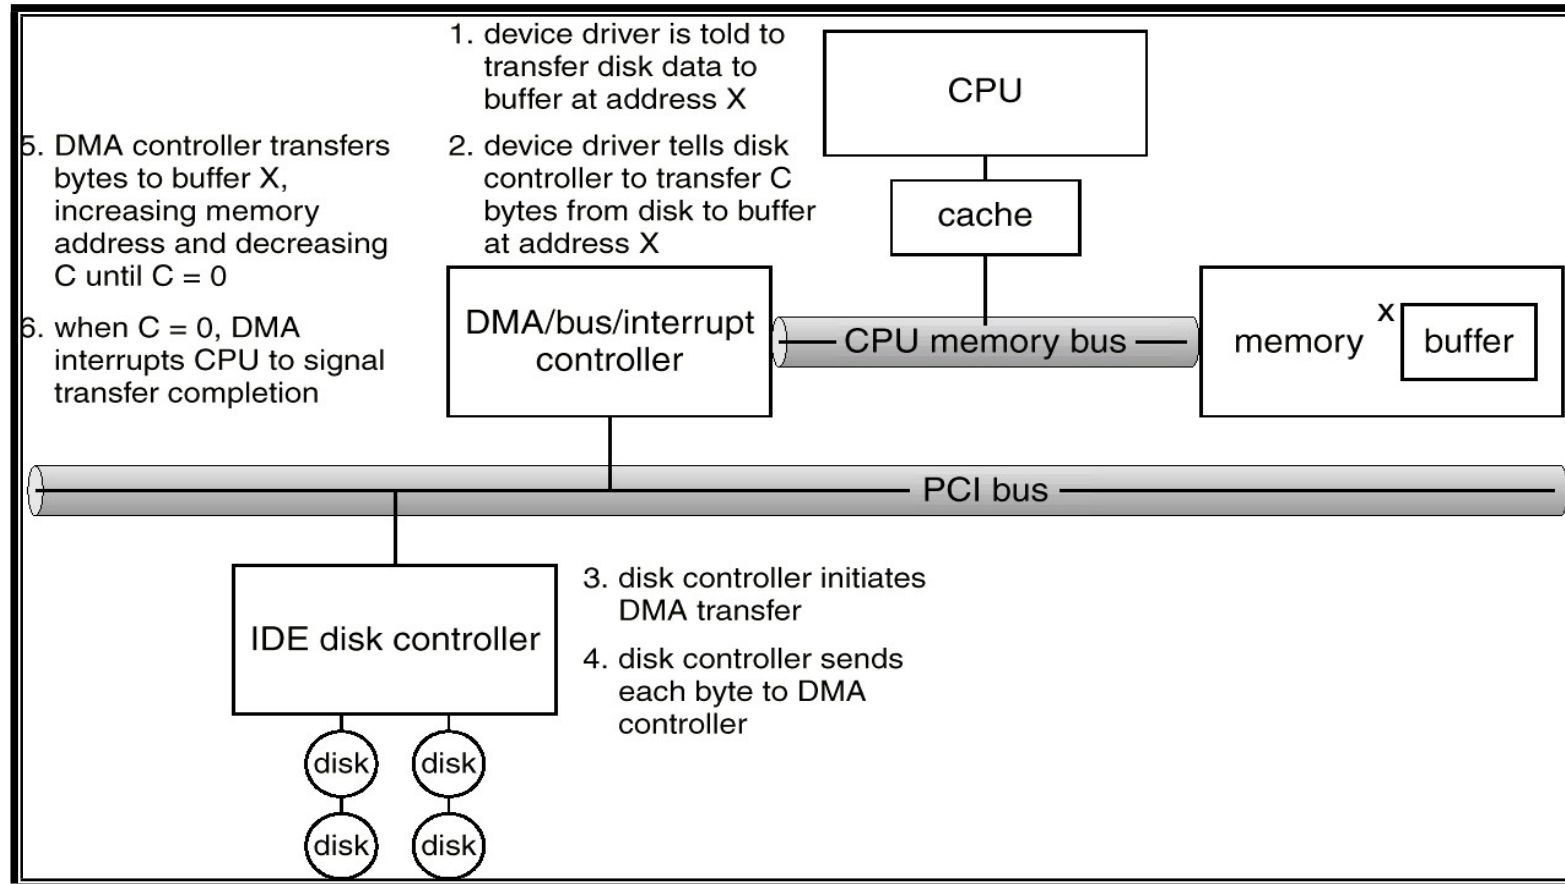
# Direct Memory Access

---

- Used to avoid programmed I/O for large data movement
  - ✓ Programmed I/O vs. DMA
- Requires DMA controller
- Bypasses CPU to transfer data directly between I/O device and memory



# Six Step Process to Perform DMA Transfer



## ■ DMA modes

### (1) Cycle stealing

- The DMA controller sneaks in and steals an occasional bus cycle from the CPU once in a while, delaying it slightly

### (2) Burst mode

- The DMA controller acquires the bus, issues a series of transfers, then releases the bus
- More efficient than cycle stealing: acquiring the bus takes time and multiple words can be transferred for the price of one bus acquisition
- It can block the CPU and other devices too long

## ■ Addressing in DMA

### (1) Physical address

- OS converts the virtual address of the intended memory buffer into a physical address and writes it into DMA controller's address register

### (2) Virtual address

- The DMA controller must use the MMU to have the virtual-to-physical translation done
  - Not common: only when the MMU is part of the memory rather than part of the CPU
- ✓ In any case, the target memory region should be pinned (not paged out) during DMA

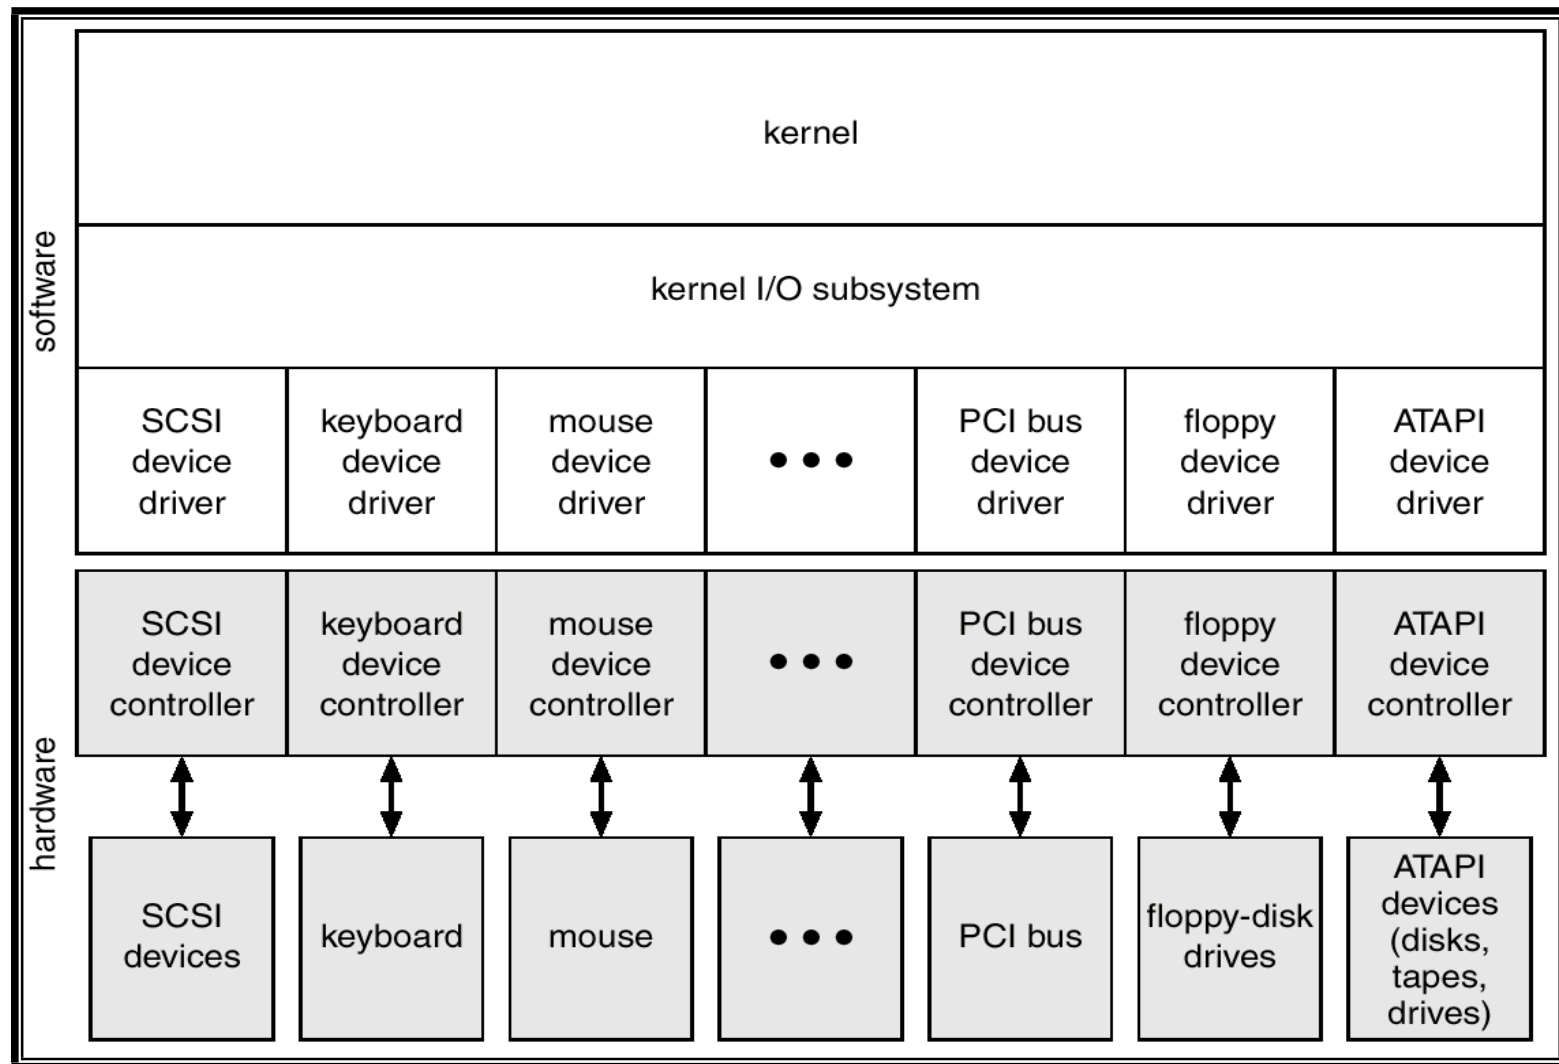
# *Application I/O Interface*

---

- I/O system calls encapsulate device behaviors in generic classes
- Device-driver layer hides differences among I/O controllers from kernel
- Devices vary in many dimensions
  - ✓ Character-stream or block
  - ✓ Sequential or random-access
  - ✓ Sharable or dedicated
  - ✓ Speed of operation
  - ✓ read-write, read only, or write only



# A Kernel I/O Structure





# Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read&write	CD-ROM graphics controller disk



# Block and Character Devices

---

- Block devices include disk drives
  - ✓ Commands include `read`, `write`, `seek`
  - ✓ Raw I/O or file-system access
  - ✓ Memory-mapped file access possible
- Character devices include keyboards, mice, serial ports
  - ✓ Commands include `get`, `put`
  - ✓ Libraries layered on top allow line editing



# Network Devices

---

- Varying enough from block and character to have own interface
- Unix and Windows NT/9i/2000 include socket interface
  - ✓ Separates network protocol from network operation
  - ✓ Includes `select` functionality
- Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)



# ***Clocks and Timers***

---

- Provide current time, elapsed time, timer
- If programmable interval time used for timings, periodic interrupts
- `ioctl` (on UNIX) covers odd aspects of I/O such as clocks and timers



# ***Blocking and Nonblocking I/O***

---

- **Blocking** - process suspended until I/O completed
  - ✓ Easy to use and understand
  - ✓ Insufficient for some needs
  
- **Nonblocking** - I/O call returns as much as available
  - ✓ User interface, data copy (buffered I/O)
  - ✓ Implemented via multi-threading
  - ✓ Returns quickly with count of bytes read or written



# Kernel I/O Subsystem

---

## ■ Scheduling


- ✓ Some I/O request ordering via per-device queue
- ✓ Some OSs try fairness

## ■ Buffering - store data in memory while transferring between devices

- ✓ To cope with device speed mismatch
- ✓ To cope with device transfer size mismatch
- ✓ To maintain “copy semantics”



# Sun Enterprise 6000 Device-Transfer Rates



# Kernel I/O Subsystem

---

- Caching - fast memory holding copy of data
  - ✓ Always just a copy
  - ✓ Key to performance
- Spooling - hold output for a device
  - ✓ If device can serve only one request at a time
  - ✓ i.e., Printing
- Device reservation - provides exclusive access to a device
  - ✓ System calls for allocation and deallocation
  - ✓ Watch out for deadlock





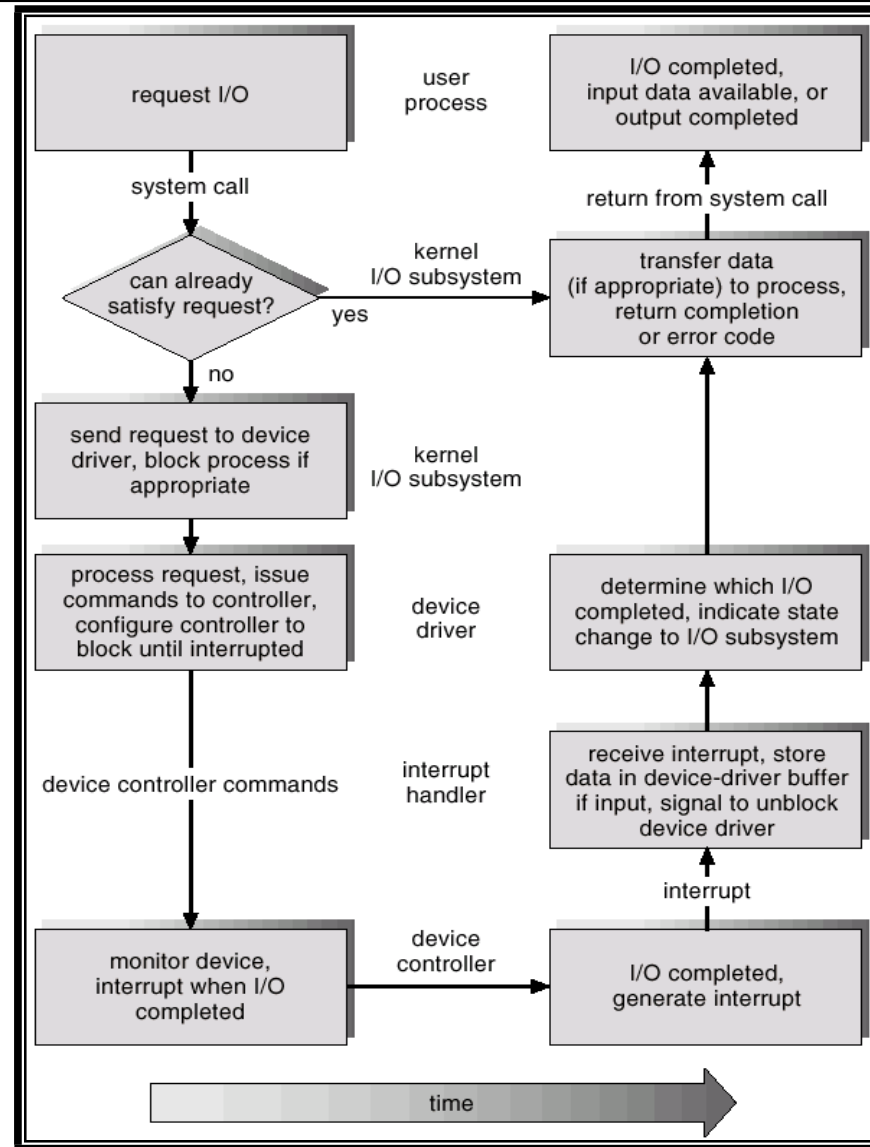
# ***Error Handling***

---

- OS can recover from disk read, device unavailable, transient write failures
- Most return an error number or code when I/O request fails
- System error logs hold problem reports



# Life Cycle of An I/O Request



# Performance

---

- I/O a major factor in system performance:
  - ✓ Demands CPU to execute device driver, kernel I/O code
  - ✓ Context switches due to interrupts
  - ✓ Data copying
  - ✓ Network traffic especially stressful



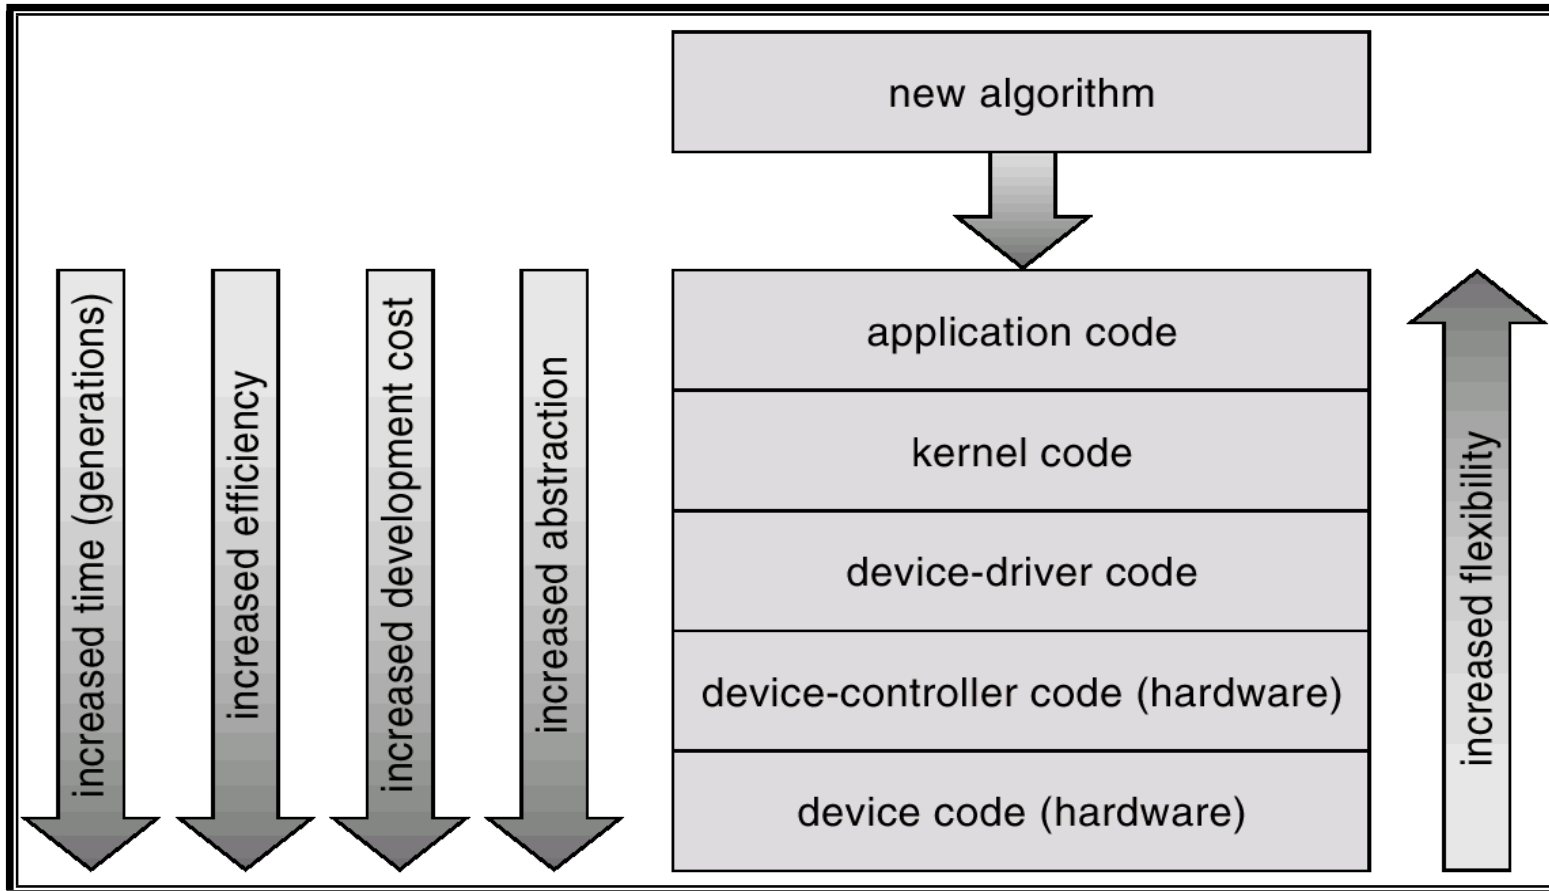
# *Improving Performance*

---

- Reduce number of context switches
- Reduce data copying
- Reduce interrupts by using large transfers, smart controllers, polling
- Use DMA
- Balance CPU, memory, bus, and I/O performance for highest throughput



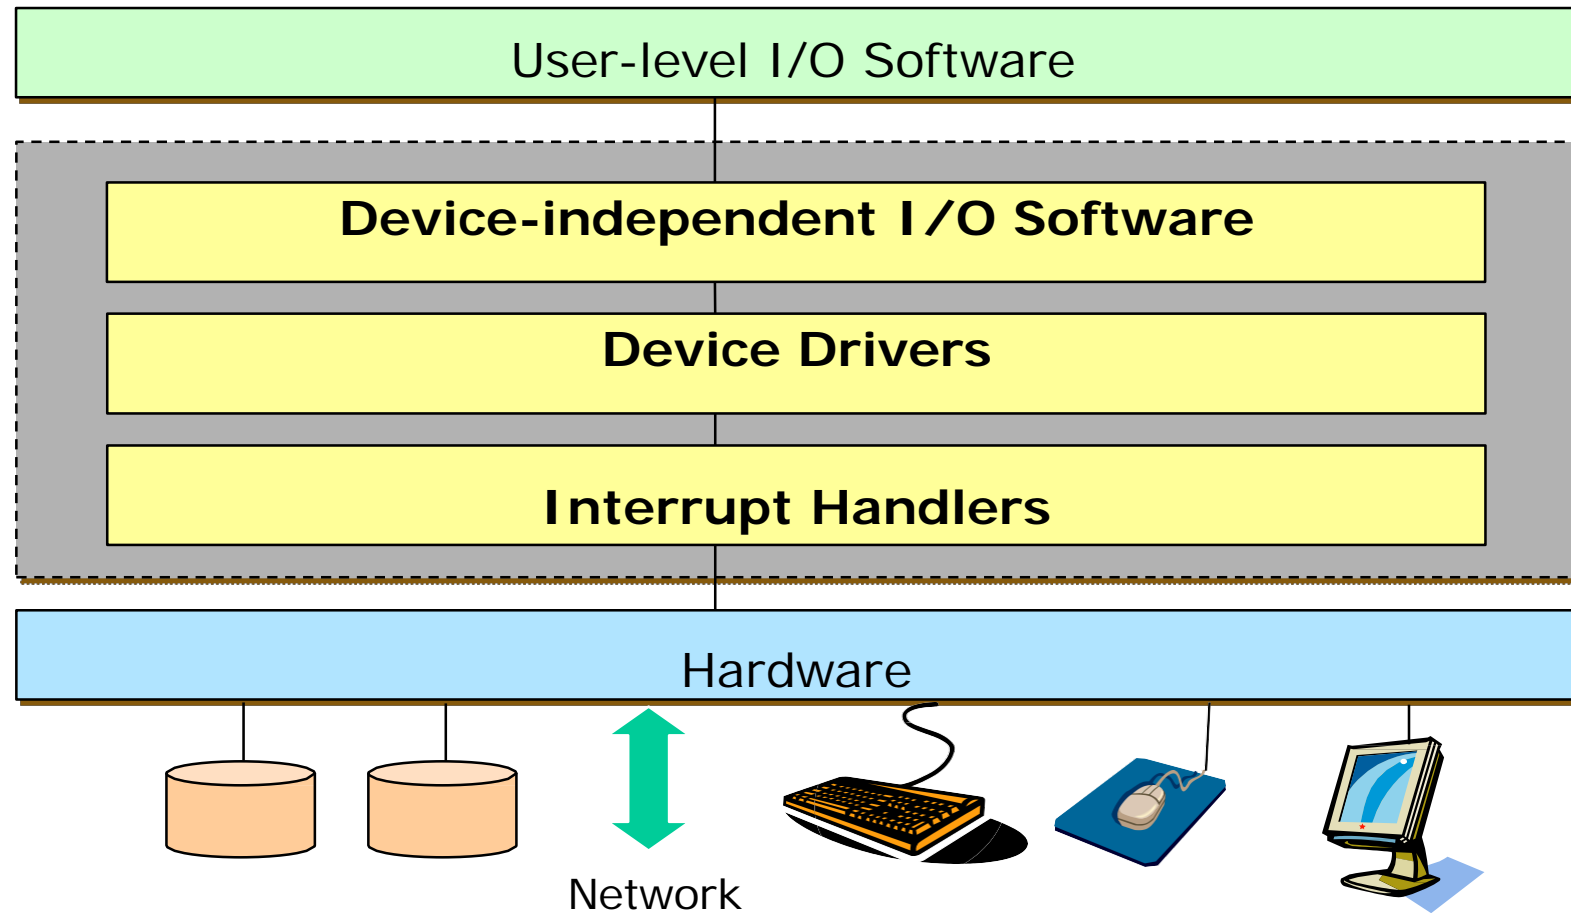
# Device-Functionality Progression




## ■ Goals

- ✓ Device independence
  - Programs can access any I/O device without specifying device in advance
- ✓ Uniform naming
  - Name of a file or device should simply be a string or an integer
- ✓ Error handling
  - Handle as close to the hardware as possible
- ✓ Synchronous vs. asynchronous
  - blocked transfers vs. interrupt-driven
- ✓ Buffering
  - Data coming off a device cannot be stored in final destination
- ✓ Sharable vs. dedicated devices
  - Disks vs. tape drives
  - Unsharable devices introduce problems such as deadlocks

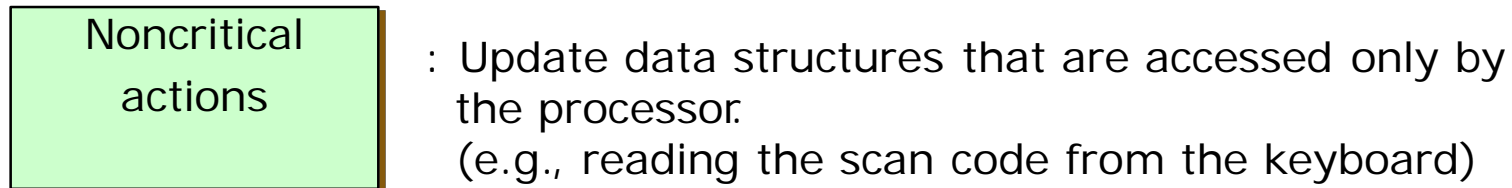
# I/O Software Layers



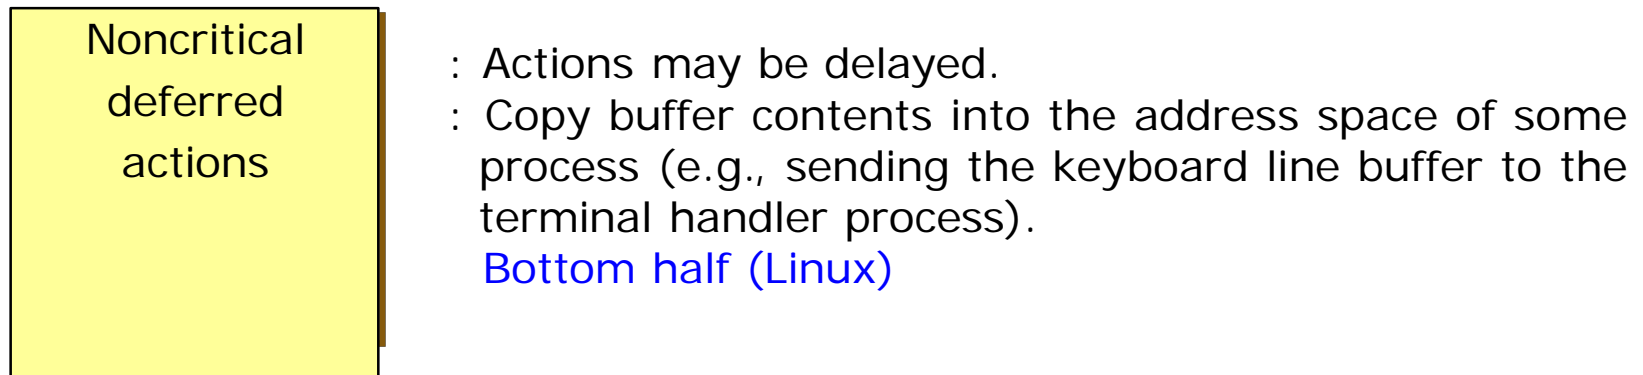
## ■ Handling interrupts



### Reenable interrupts



### Return from interrupts





## ■ Device drivers

- ✓ Device-specific code to control each I/O device interacting with device-independent I/O software and interrupt handlers
- ✓ Requires to define a well-defined model and a standard interface of how they interact with the rest of the OS
- ✓ Implementing device drivers:
  - Statically linked with the kernel
  - Selectively loaded into the system during boot time
  - Dynamically loaded into the system during execution (especially for hot pluggable devices)

# Device-Independent I/O Software

---

## ■ Uniform interfacing for device drivers

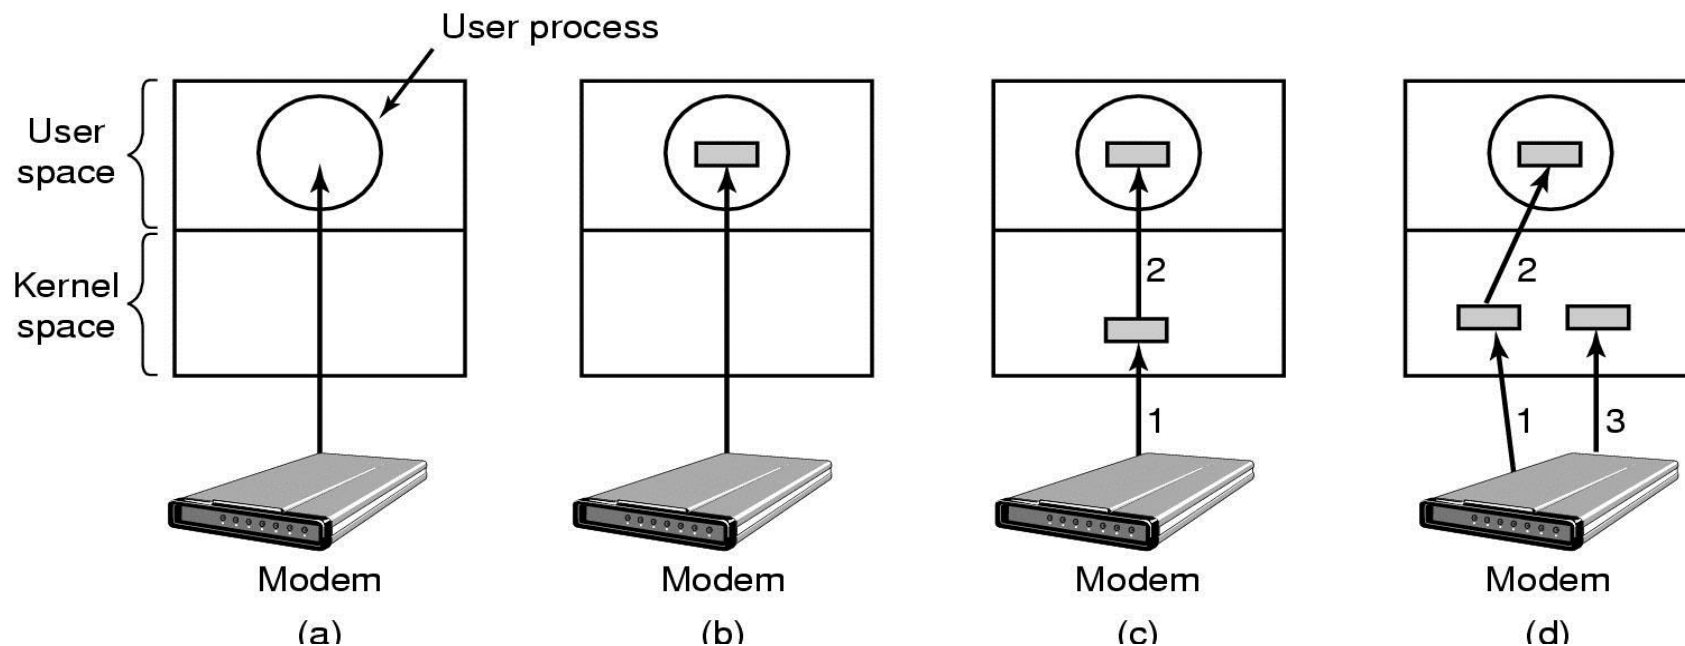
- ✓ In Unix, devices are modeled as special files
  - They are accessed through the use of system calls such as open(), read(), write(), close(), ioctl(), etc.
  - A file name is associated with each device
- ✓ Major device number locates the appropriate driver
  - Minor device number (stored in i-node) is passed as a parameter to the driver in order to specify the unit to be read or written
- ✓ The usual protection rules for files also apply to I/O devices



# Device-Independent I/O Software (Cont'd)

## ■ Buffering

- ✓ (a) Unbuffered
- ✓ (b) Buffered in user space
- ✓ (c) Buffered in the kernel space
- ✓ (d) Double buffering in the kernel



# *Device-Independent I/O Software (Cont'd)*

---

## ■ Error reporting

- ✓ Many errors are device-specific and must be handled by the appropriate driver, but the framework for error handling is device independent
- ✓ Programming errors vs. actual I/O errors
- ✓ Handling errors
  - Returning the system call with an error code
  - Retrying a certain number of times
  - Ignoring the error
  - Killing the calling process
  - Terminating the system



# Device-Independent I/O Software (Cont'd)

---

## ■ Allocating and releasing dedicated devices

- ✓ Some devices cannot be shared

- (1) Require processes to perform `open()`'s on the special files for devices directly

  - The process retries if `open()` fails

- (2) Have special mechanisms for requesting and releasing dedicated devices

  - An attempt to acquire a device that is not available blocks the caller

## ■ Device-independent block size

- ✓ Treat several sectors as a single logical block

- ✓ The higher layers only deal with abstract devices that all use the same block size



## ■ Provided as a library

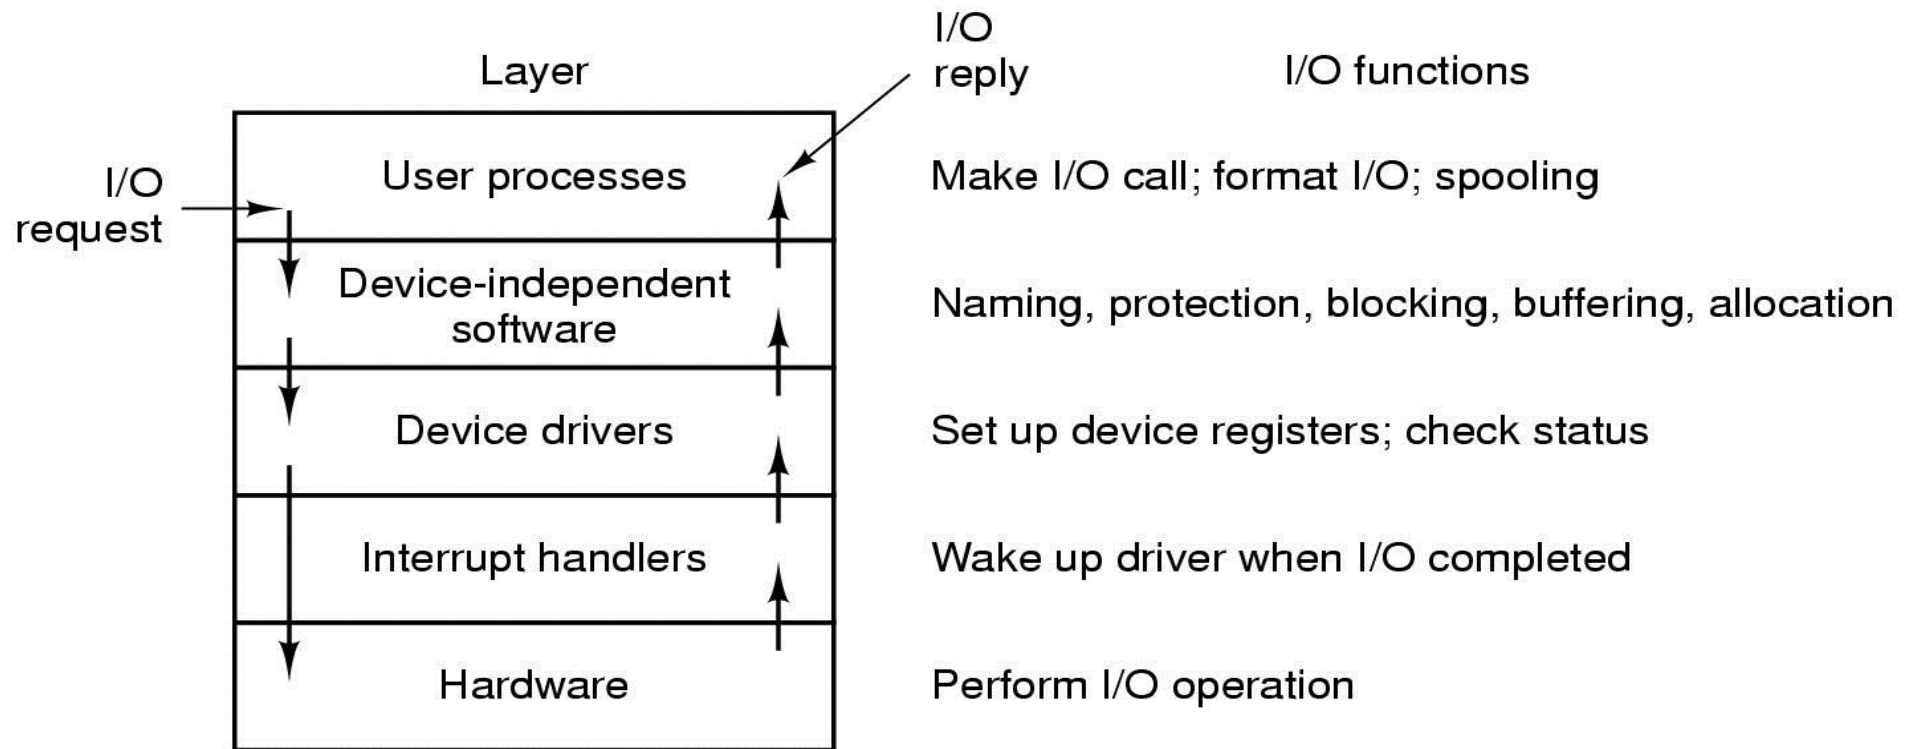
- ✓ Standard I/O library in C
  - `fopen()` vs. `open()`

## ■ Spooling

- ✓ A way of dealing with dedicated I/O devices in a multiprogramming system
- ✓ Implemented by a daemon and a spooling directory
- ✓ Printers, network file transfers, USENET news, mails, etc.



# I/O Systems Layers



## ■ I/O hardware

- ✓ I/O instructions control I/O devices
  - Direct I/O vs. Memory-mapped I/O
- ✓ I/O events notification
  - Polled I/O vs. Interrupt-driven I/O
- ✓ DMA (Direct Memory Access)

## ■ I/O software

- ✓ Application I/O interface: I/O system call
  - open, read, write, close, ioctl, ...
- ✓ Kernel I/O structure
  - Device-independent I/O subsystem
  - Device drivers (device specific)
  - Interrupt handlers (device specific)



## ■ Performance issues

- ✓ Reduce number of context switches
- ✓ Reduce data copying
- ✓ Reduce interrupts
- ✓ Use DMA
- ✓ Balance CPU, memory, bus, and I/O performance for higher system-level throughput

